

1 Modular Program Slicing Through Ownership

2 ANONYMOUS AUTHOR(S)

3 Program slicing, or identifying the subset of a program relevant to a value, relies on understanding the
 4 dataflow of a program. In languages with functions and mutable pointers like C or Java, tracking dataflow
 5 has historically required whole-program analysis, which can be slow and challenging to integrate in
 6 practice. Advances in type systems have shown how to modularly track dataflow through the concept of
 7 ownership. We demonstrate that ownership can modularize program slicing by using types to compute a
 8 provably sound and reasonably precise approximation of mutation. We design a modular slicing algorithm for
 9 Rust, an industrial-grade ownership-based programming language. We prove the algorithm's soundness as a
 10 form of noninterference on the Oxide formal model of Rust. Then we empirically evaluate the precision of the
 11 slicer, showing that modular slices are the same as whole-program slices in 95% of cases drawn from large
 12 Rust codebases.

13 Additional Key Words and Phrases: program slicing, ownership, symbolic procedure summaries, modular
 14 static analysis, rust

17 1 INTRODUCTION

18 Program slicing is the task of identifying the subset of a program relevant to computing a value of
 19 interest. The concept of slicing was introduced 40 years ago when Weiser [1982] demonstrated that
 20 programmers mentally construct slices while debugging. Since then, hundreds of papers have been
 21 published on implementing automated program slice, as surveyed by Silva [2012]; Xu et al. [2005].
 22 Despite these efforts, a review of slicers found “slicing-based debugging techniques are rarely used
 23 in practice” [Parnin and Orso 2011].¹

24 A major challenge for slicing is addressing the underlying program analysis problems. At a
 25 high level, slicing is about dataflow — if x is relevant, then any means by which data flows into
 26 x are also relevant. In today’s programming languages, analyzing dataflow is difficult because of
 27 the interaction of two features: functions and pointers. For example, imagine slicing a value in a
 28 function f which calls a function g . In a language without side-effects, then the only relevance g
 29 could possibly have in f is its return value. But in a language that allows effects such as mutation
 30 on pointers, g could modify data used within f , requiring a pointer analysis. Moreover, if f is a
 31 higher-order function parameterized on g , then the slice must consider all the possible functions
 32 that g could be, i.e. control-flow analysis.

33 The standard solution for analyzing programs with pointers and functions is *whole-program*
 34 *analysis*. That is, for a given function of interest, analyze the definitions of all of the function’s
 35 callers and callees in the current codebase. However, whole-program analysis suffers from a few
 36 logistical and conceptual issues:

- 37 (1) *Analysis time scales with the size of the whole program*: the time complexity of whole-program
 38 analysis scales either polynomially or exponentially with the number of call sites in the program,
 39 depending on context-sensitivity [Might et al. 2010]. In practice, this means more complex
 40 codebases can take substantially longer to analyze. For instance, the recent PSEGPT pointer
 41 analysis tool [Zhao et al. 2018] takes 1 second on a codebase of 282,000 lines of code and 3
 42 minutes on a codebase of 2.2 million lines of code.

43
 44 ¹The only open-source, functioning slicers the authors could find are Frama-C [Cuoq et al. 2012] and dg [Chalupa 2016].
 45 Slicing tools for Java like Kaveri [Jayaraman et al. 2005] no longer work. The most industrial-strength slicing tool,
 46 CodeSurfer [Balakrishnan et al. 2005] was GrammaTech’s proprietary technology and appears to no longer exist.

47 2018. 2475-1421/2018/1-ART1 \$15.00

48 <https://doi.org/>

- 50 (2) *Analysis requires access to source code for the whole program*: an assumption of analyzing a whole
 51 program is that a whole program is actually accessible. However, many programs use libraries
 52 that are shipped as pre-compiled objects with no source code, either for reasons of efficiency or
 53 intellectual property.
 54 (3) *Analysis results are anti-modular*: when analyzing a particular function, relying on calling
 55 contexts to analyze the function's inputs means that any results are not universal. Calling-
 56 context-sensitive analysis determine whether two pointers alias *in the context of the broader*
 57 *codebase*, so alias analysis results can change due to modifications in code far away from the
 58 current module.

59 These issues are not new – Rountev et al. [1999] and Cousot and Cousot [2002] observed the
 60 same two decades ago when arguing for modular static analysis. The key insight arising from their
 61 research is that static analysis can be modularized by computing *symbolic procedure summaries*.
 62 For instance, Yorsh et al. [2008] show how to automatically summarize which inputs and outputs
 63 are possibly null for a given Java function. The analysis is modular because a function's summary
 64 can be computed only given the summaries, and not definitions, of callees in the function. In such
 65 prior work, the language of symbolic procedure summaries has been defined in a separate formal
 66 system from the programming language being analyzed, such as the micro-transformer framework
 67 of Yorsh et al. [2008].

68 Our work begins with the observation: *function type signatures are symbolic procedure summaries*.
 69 The more expressive a language's type system, the more behavior that can be summarized by a
 70 type. Nearly all work on program slicing, dataflow analysis, and procedure summaries has operated
 71 on C, Java, or equivalents. These languages have impoverished type systems, and so any interesting
 72 static analysis requires a standalone abstract interpreter. However, if a language's type system were
 73 expressive enough to encode information about dataflow, then a function's type signature could be
 74 used to reason about the aliasing and side effects needed for slicing. Moreover, a function's type
 75 signature is required information for a compiler to export when building a library. Using the type
 76 system for dataflow analysis therefore obviates the logistical challenge of integrating an external
 77 analysis tool into a complex build system.

78 Today, the primary technique for managing dataflow with types is *ownership*. Ownership is a con-
 79 cept that has emerged from several intersecting lines of research on linear logic [Girard 1987], class-
 80 based alias management [Clarke et al. 1998], and region-based memory management [Grossman
 81 et al. 2002]. Generally, ownership refers to a system where values are owned by an entity, which can
 82 temporarily or permanently transfer ownership to other entities. The type system then statically
 83 tracks the flow of ownership between entities. Ownership-based type systems enforce the invariant
 84 that values are not simultaneously aliased and mutated, either for the purposes of avoiding memory
 85 errors, data races, or abstraction violations.

86 Our thesis is that ownership can modularize program slicing by using types to compute a provably
 87 sound and reasonably precise approximation of the necessary dataflow information. We build this
 88 thesis in five parts:

- 89 (1) We provide an intuition for the relationship between ownership and slicing by describing
 90 how ownership works in Rust, an industrial-grade ownership-based programming language
 91 (Section 2).
 92 (2) We formalize an algorithm for modular slicing as an extension to the type system of Ox-
 93 ide [Weiss et al. 2019], a formal model of Rust's static and dynamic semantics (Section 3 and
 94 Section 4).
 95 (3) We prove the soundness of this algorithm as a form of noninterference, building on the
 96 connection between slicing and information flow established by Abadi et al. [1999] (proof

sketch in Section 4.2 and full proof in Appendix A.2, located in supplementary material for submission).
 101 (4) We describe an implementation of the slicing algorithm for Rust, translating the core insights
 102 of the algorithm to work on a lower-level control-flow graph (Section 5).
 103 (5) We evaluate the precision of the modular Rust slicer against a whole-program slicer on a
 104 dataset of 10 codebases with a total of 280k LOC. We find that modular slices are the same size
 105 as whole-program slices 95.4% of the time, and are on average 7.6% larger in the remaining
 106 4.6% of cases (Section 6).

107

2 PRINCIPLES

A backwards static slice is the subset of a program that could influence a particular value (backwards) under any possible execution (static). A slice is defined with respect to a slicing criterion, which is a variable at a particular point in a program. In this section, we provide an intuition for how slices interact with different features of the Rust programming language, namely: places (2.1), references (2.2), function calls (2.3), and interior mutability (2.4).

114

2.1 Places

A place is a reference to a concrete piece of data in memory, like a variable `x` or path into a data structure `x.field`. Slices on places are defined by bindings, mutation, and control flow. For instance, the Rust snippet on the right shows the slice in orange of a place in green. The assignment `x = y` means `y` is relevant for the slice, so the statement `let y = 2` is relevant as well. Because `z` is not used in the computation of `x`, then `let z = 3` is not relevant. Additionally, because `x = y` overwrites the previous value of `x`, then the original assignment `x = 1` is not relevant either.

```
1  let mut x = 1;
2  let mut y = 2;
3  if y > 0 { x = 3; }
4  else      { y = 4; }
5  println!("{} {}", x, y);
```

mutation conflicts with a particular path into the data structure. For example, consider slicing on a tuple as in the three snippets below (note that `t.n` gets the *n*-th field of the tuple `t`):

<pre>1 let mut t = (0, 1, 2); 2 t = (3, 4, 5); 3 t.0 = 6; 4 t.1 = 7; 5 println!("{}?", t);</pre>	<pre>1 let mut t = (0, 1, 2); 2 t = (3, 4, 5); 3 t.0 = 6; 4 t.1 = 7; 5 println!("{}?", t.0);</pre>	<pre>1 let mut t = (0, 1, 2); 2 t = (3, 4, 5); 3 t.0 = 6; 4 t.1 = 7; 5 println!("{}?", t.2);</pre>
---	---	---

In this program, when slicing on `t`, changing the value of a field of a structure changes the value of the whole structure, so `t.1 = 7` is part of the slice on `t`. However, when slicing on `t.0`, the path `t.0` is disjoint from the path `t.1`, so `t.1 = 7` is not part of the slice on `t.0`. Similarly, when slicing on `t.2`, the only relevant assignment is `t = (3, 4, 5)`. More generally, a place conflicts with another place if either's path is a prefix of the other's. For instance, `t.0` conflicts with both `t` (parent) and `t.0.1` (child) but not `t.1` (sibling). This leads to the first slicing principle:

144

Principle 1 (Slicing principle for places). *A mutation to a place is a mutation to all conflicting places.*

147

```
1  let mut x = 1;
2  let y = 2;
3  let z = 3;
4  x = y;
5  println!("{} {}", x, z);
```

148 This principle provides an intuition for making an algorithm that constructs slices. For instance,
 149 take the last example above on the left. On line 4, when `t.1` is mutated, that mutation is registered
 150 as part of the slice on every conflicting place, specifically `t` and `t.1`.

153 2.2 References

154 Pointers are the first major challenge for slicing. A mutation to a dereferenced pointer is a mutation
 155 to any place that is possibly pointed-to, so such places must be known to the slicer. For example:

156 Rust has two distinct types of pointers, which are called “references”
 157 to distinguish them from “raw pointers” with C-like behavior (discussed
 158 in Section 2.4). For a given type `T`, there are immutable references of type
 159 `&T`, and mutable references of type `&mut T` which correspond respectively
 160 to the expressions `&x` and `&mut x`. Because `y` points to `x`, then the mutation
 161 through `y` is relevant to the read of `*z`. We refer to the left-hand side of
 162 assignment statements like `*y` as “place expressions”, since they could include dereferences.

```
1 let mut x = 1;
2 let y = &mut x;
3 *y = 2;
4 let z = &x;
5 println!("{}", *z);
```

163 The task of determining what a reference can point-to is called *pointer analysis*. While many
 164 methods exist for pointer analysis [Smaragdakis and Balatsouras 2015], our first key insight is
 165 that Rust’s ownership types implicitly perform a kind of modular pointer analysis that we can
 166 leverage for slicing. To understand why, we first need to describe two ingredients: the goal, i.e.
 167 what ownership is trying to accomplish, and the mechanism, i.e. how ownership-checking is
 168 implemented in the type system.

169 The core goal of ownership is eliminating simultaneous aliasing and mutation. In Rust, achieving
 170 this goal enables the use of references without garbage collection while retaining memory safety.
 171 For instance, these three classes of errors are all caught at compile-time:

<pre>1 // Dangling reference 2 let p = { 3 let x = 1; &x 4 }; 5 let y = *p;</pre>	<pre>1 // Use-after-free 2 let d = tempdir(); 3 let d2 = &d; 4 d.close(); 5 let p = d2.path();</pre>	<pre>1 // Iterator invalidation 2 let mut v = vec![1,2]; 3 for x in v.iter() { 4 v.push(*x); 5 }</pre>
---	--	--

172 From left-to-right: the dangling references is caught because `x` is deallocated at the end of
 173 scope on line 4, which is a mutation, conflicting with the alias `&x`. The use-after-free is caught
 174 because `d.close()` requires ownership of `d`, which prevents an alias `d2` from being live. The iterator
 175 invalidation case is subtler: `x` is a pointer to data within `v`. However, `v.push(*x)` could resize `v`
 176 which would copy/deallocate all vector elements to a new heap location, invalidating all pointers
 177 to `v`. Hence `v.push(*x)` is a simultaneous mutation and alias of the vector.

178 Catching these errors requires understanding which places are pointed by which references.
 179 For instance, knowing that `x` points to an element of `v` and not just any arbitrary `i32`. The key
 180 mechanism behind these ownership checks is *lifetimes*.

```
1 let mut x: i32 = 1;
2 let y: &'1 i32 = &'0 mut x;
3 *y = 2;
4 let z: &'3 i32 = &'2 x;
5 println!("{}", *z);
```

181 Each reference expression and type has a corresponding
 182 lifetime, written explicitly in the syntax '`n`' on the left, where `n`
 183 is an arbitrary and unique number. The name “lifetime” implies
 184 a model of lifetimes as the live range of the reference. Prior
 185 work on region-based memory management like [Tofte and Talpin \[1997\]](#) and [Grossman et al. \[2002\]](#) use this model.

186 However, recent work from [Matsakis \[2018\]](#) and [Weiss et al. \[2019\]](#) have devised an alternative
 187 model of lifetimes as “provenances” or “origins” that more directly correspond to a pointer analysis.
 188 In essence, a lifetime is the set of places that a reference could point-to. For the above example,
 189

that would be ' $n = \{x\}$ ' for all n , because each reference points to x . As a more interesting example, consider the code on the left.

```

1  let mut x = 1;
2  let mut y = 2;
3  let z: &'2 mut i32 = if true {
4      &'0 mut x
5  } else {
6      &'1 mut y
7  };
8  let w: &'4 mut i32 = &'3 mut *z;
9  *w = 1;

```

There, lifetimes for borrow expressions are assigned to the place being borrowed, so ' $\theta = \{x\}$ ' and ' $\iota = \{y\}$ '. Because z could be assigned to either reference, then ' $\tau = \theta \cup \iota = \{x, y\}$ '. An expression of the form $\&p$ is called a "reborrow", as the underlying address is being passed from one reference to another. To register that a reference is reborrowed, the reborrowed place is also added to the lifetime, so ' $\kappa = \tau \cup \{\star z\}$ '. More generally:

Principle 2 (Slicing principle for references). *The lifetime of a reference contains all potential aliases of what the reference points-to.*

In the context of slicing, then to determine which places could be modified by a particular assignment, one only needs to look up the aliases in the lifetime of references. For instance, $*w = 1$ would be part of a slice on $\star z$, because $\star z$ is in the lifetime ' τ ' of w .

2.3 Function calls

The other major challenge for slicing is function calls. For instance, consider slicing a call to an arbitrary function f with various kinds of inputs.²

```

1  let x = String::from("x");
2  let y = String::from("y");
3  let mut z = String::from("z");
4  let w = f(x, &y, &mut z);
5  println!("{} {} {}", y, z, w);

```

The standard approach to slicing f would be to inspect the definition of f , and recursively slice it by translating the slicing criteria from caller to callee (e.g. see [Weiser \[1982\]](#) for an example). However, our goal is to avoid using the definition of f (i.e. a whole-program analysis) for the reasons described in Section 1.

To modularly slice through function calls, we need to approximate the effects of f in a manner that is sound, but also as precise as possible. Put another way, what mutations could possibly occur as a result of calling f ? Consider the three cases that arise in the code above.

- Passing a value x of type `String` (or generally of type T) moves the value into f . Therefore it is an ownership error to refer to x after calling f and we do not need to consider slices on x after f .
- Passing a value y of type `&String` (or $\&T$) passes an immutable reference. Immutable references cannot be mutated, therefore y cannot change in f .³
- Passing a value z of type `&mut String` (or $\&mut T$) passes a mutable reference, which could possibly be mutated. This case is therefore the only observable effect f apart from its return value.

Without inspecting f , we cannot know how a mutable reference is modified, so we have to conservatively assume that every argument was used as input to a mutation. Therefore the modular slice of each variable looks as in the snippets below:

²Why is `String::from` needed? The literal "Hello world" has type `&'static str`, meaning an immutable reference to the binary's string pool which lives forever. The function `String::from` converts the immutable reference into a value of type `String`, which stores its contents on the heap and allows the string to be mutated.

³A notable detail to the safety of immutable references is that immutability is transitive. For instance, if $b = \&mut a$ and $c = \&b$, then a is guaranteed not to be mutated through c . This stands in contrast to other languages with pointers like C and C++ where the `const` keyword only protects values from mutation at the top-level, and not into the interior fields.

```

246 1 let x = String::from("x");
247 2 let y = String::from("y");
248 3 let mut z = String::from("z");
249 4 let w = f(x, &y, &mut z);
250 5 println!("{} {} {}", y, z, w);
251

```

```

252 1 let x = String::from("x");
253 2 let y = String::from("y");
254 3 let mut z = String::from("z");
255 4 let w = f(x, &y, &mut z);
256 5 println!("{} {} {}", y, z, w);
257

```

was not an input!). Similarly, functions could potentially read arbitrary data (e.g. global variables) that would influence mutations apart from just the arguments.

However, allowing such pointer manipulation would easily break ownership safety, since fundamentally it permits unchecked aliasing. Hence, our principle:

Principle 3 (Slicing principle for function calls). *When calling a function, (a) only mutable references in the arguments can be mutated, and (b) the mutations and return value are only influenced by the arguments.*

This principle is essentially a worst-case approximation to the function's effects. It is the core of how we can modularly slice programs, because a function's definition does not have to be inspected to analyze what it can mutate.

A caveat to this principle is global variables: (Principle 3-a) is not true with mutable globals, and (Principle 3-b) is not true with read-only globals. Mutable globals are disallowed by the rules of ownership, as they are implicitly aliased and hence disallowed from being mutable. However, read-only globals are ownership-safe (and hence permitted in Rust). For simplicity we do not consider read-only globals in this work.

Another notable detail is the interaction of function calls and lifetimes. Pointer analysis, like slicing, has historically been done via whole-program analysis for maximum precision. However, Rust can analyze lifetimes (and subsequently what references point-to) modularly just by looking at the type signature of a called function using *lifetime parameters*. Consider the function `Vec::get_mut` that returns a mutable reference to an element of a vector. For instance, `vec![5, 6].get_mut(0)` returns a mutable reference to the value 5. This function has the type signature:

```
Vec::get_mut : ∀ 'a, T . (&'a mut Vec<T>, usize) -> &'a mut T
```

Because this type signature is parametric in the lifetime '`'a`', it can express the constraint that the output reference `&'a mut T` must have the same lifetime as the input reference `&'a mut Vec<T>`. Therefore the returned pointer is known to point to the same data as the input pointer, but without inspecting the definition of `get_mut`.

2.4 Interior mutability

The previous sections describe a slicing strategy for the subset of Rust known as “safe Rust”, that is programs which strictly adhere to the rules of ownership. Importantly, Rust also has the `unsafe` feature that gives users access to raw pointers, or pointers with similar unchecked behavior to C. Most commonly, `unsafe` code is used to implement APIs that satisfy ownership, but not in a manner that is deducible by the type system. For example, shared mutable state between threads:

```

1 let x = String::from("x");
2 let y = String::from("y");
3 let mut z = String::from("z");
4 let w = f(x, &y, &mut z);
5 println!("{} {} {}", y, z, w);

```

Note that like `z` (above), the return value `w` (left) is also assumed to be influenced by every input to `f`. Implicit in these slices are additional assumptions about the limitations of `f`. For example, in C, a function could manufacture a pointer to the stack frame above it and mutate the values, meaning `f` could mutate `y` (even if `y`

was not an input!). Similarly, functions could potentially read arbitrary data (e.g. global variables) that would influence mutations apart from just the arguments.

However, allowing such pointer manipulation would easily break ownership safety, since fundamentally it permits unchecked aliasing. Hence, our principle:

Principle 3 (Slicing principle for function calls). *When calling a function, (a) only mutable references in the arguments can be mutated, and (b) the mutations and return value are only influenced by the arguments.*

This principle is essentially a worst-case approximation to the function's effects. It is the core of how we can modularly slice programs, because a function's definition does not have to be inspected to analyze what it can mutate.

A caveat to this principle is global variables: (Principle 3-a) is not true with mutable globals, and (Principle 3-b) is not true with read-only globals. Mutable globals are disallowed by the rules of ownership, as they are implicitly aliased and hence disallowed from being mutable. However, read-only globals are ownership-safe (and hence permitted in Rust). For simplicity we do not consider read-only globals in this work.

Another notable detail is the interaction of function calls and lifetimes. Pointer analysis, like slicing, has historically been done via whole-program analysis for maximum precision. However, Rust can analyze lifetimes (and subsequently what references point-to) modularly just by looking at the type signature of a called function using *lifetime parameters*. Consider the function `Vec::get_mut` that returns a mutable reference to an element of a vector. For instance, `vec![5, 6].get_mut(0)` returns a mutable reference to the value 5. This function has the type signature:

```
Vec::get_mut : ∀ 'a, T . (&'a mut Vec<T>, usize) -> &'a mut T
```

Because this type signature is parametric in the lifetime '`'a`', it can express the constraint that the output reference `&'a mut T` must have the same lifetime as the input reference `&'a mut Vec<T>`. Therefore the returned pointer is known to point to the same data as the input pointer, but without inspecting the definition of `get_mut`.

2.4 Interior mutability

The previous sections describe a slicing strategy for the subset of Rust known as “safe Rust”, that is programs which strictly adhere to the rules of ownership. Importantly, Rust also has the `unsafe` feature that gives users access to raw pointers, or pointers with similar unchecked behavior to C. Most commonly, `unsafe` code is used to implement APIs that satisfy ownership, but not in a manner that is deducible by the type system. For example, shared mutable state between threads:

```

295 1 let value = Arc::new(Mutex::new(0));
296 2
297 3 let value_ref = value.clone();
298 4 thread::spawn(move || {
299 5     *value_ref.lock().unwrap() += 1;
300 6 }).join().unwrap();
301 7
302 8 assert!(*value.lock().unwrap() == 1);
303

```

The mutex is ownership-safe only because its implementation ensures that both threads cannot simultaneously access the underlying value in accordance with the system mutex's semantics. For our purposes, the aliasing between `value` and `value_ref` is not possible to observe using the type system alone. For example, in our algorithm, slicing on `value` would *not* include mutations to `value_ref`. This is because the data inside the mutex has type `*mut i32` (a raw pointer), and without a lifetime attached, our algorithm has no way to determine whether `value` and `value_ref` are aliases just by inspecting their types.

More broadly, modular slicing is only sound for safe Rust. The point of this work is to say: when a program can be statically determined to satisfy the rules of ownership, then modular slicing is sound. The principles above help clarify the specific assumptions made possible by ownership, which are otherwise impossible to make in languages like C or Java. Astrauskas et al. [2020] found that 76.4% of published Rust projects contain no unsafe code, suggesting that safe Rust is more common than not. However, their study does not account for safe Rust built on internally-unsafe abstractions like Mutex, so it is difficult to estimate the true likelihood of soundness in practice. We discuss the issue of slicing with unsafe code further in Section 6.4.1.

3 FORMAL MODEL

To build an algorithm from these principles, we first need a formal model to describe and reason about the underlying language. Rather than devise our own, we build on the work of Weiss et al. [2019]: Oxide is a model of (safe) Rust's surface language with a formal static and dynamic semantics, along with a proof of syntactic type soundness. Importantly, Oxide uses a provenance model of lifetimes which we leverage for our slicing algorithm.

We will incrementally introduce the aspects of Oxide's syntax and semantics as necessary to understand our principles and algorithm. We describe Oxide's syntax (3.1), static semantics (3.2) and dynamic semantics (3.3), and then apply these concepts to formalize the slicing principles of the previous section (3.4).

3.1 Syntax

Figure 1 shows a subset of Oxide's syntax along with a labeled example. An Oxide program consists of a set of functions Σ (the “global environment”), where each function body is an expression e . The syntax is largely the same as Rust's with a few exceptions:

- Lifetimes are called “provenances”, and they are both explicit in expressions and types throughout the program, and initially bound via `letprov` expressions or as function parameters.
- Rather than having immutable references `&'a τ` and mutable references `&'a mut τ`, Oxide calls them “shared” references `&ρ shrd τ` and “unique” references `&ρ uniq τ`.

In this snippet, two threads have ownership over two values of type `Arc<Mutex<i32>>` which internally point to the same number. Both threads can call `Mutex::lock` which takes an immutable reference to an `&Mutex<i32>` and returns a mutable reference `&mut i32` to the data inside.⁴ This nominally violates ownership, as the data is aliased (shared by two threads) and mutable (both can mutate).

⁴Technically the returned type is a `LockResult<MutexGuard<'a, i32>>` but the distinction isn't relevant here.

```

344 Path  $q ::= \varepsilon \mid n.q$  Provenances  $\rho ::= \varrho \mid r$ 
345 Places  $\pi ::= x.q$  Ownership Qualifiers  $\omega ::= \text{shrd} \mid \text{uniq}$ 
346
347 Place Expressions  $p ::= x \mid *p \mid p.n$  Base Types  $\tau^B ::= \text{unit} \mid \text{u32} \mid \text{bool}$ 
348 Constants  $c ::= () \mid n \mid \text{true} \mid \text{false}$  Sized Types  $\tau^{\text{SI}} ::= \tau^B \mid \&\rho \omega \tau^{\text{XI}} \mid (\tau_1^{\text{SI}}, \dots, \tau_n^{\text{SI}})$ 
349
350 Expressions  $e ::= c \mid p \mid \&r \omega p \mid \text{letprov}\langle r \rangle \{ e \} \mid \text{if } e_1 \text{ then } e_2 \text{ else } e_3$ 
351
352 Global Entries  $\varepsilon ::= \text{fn } f\langle \bar{\psi}, \bar{\varrho}, \bar{\alpha} \rangle(x : \tau_a^{\text{SI}}) \rightarrow \tau_r^{\text{SI}} \text{ where } \bar{\varrho}_1 : \bar{\varrho}_2 \{ e \}$ 
353
354 Global Environment  $\Sigma ::= \bullet \mid \Sigma, \varepsilon$ 
355
356 (a) Subset of Oxide syntax, reproduced from Weiss et al. [2019, p. 8]. The only difference in this subset is that
357 closures are eliminated and functions are simplified to take one argument.
358
359
360 Variable  $x$  Sized Type  $\tau^{\text{SI}}$  Expression  $e$  Place  $\pi$ 
361
362
363
364
365
366
367
368
369
370 (b) Syntactic forms and corresponding metavariables labeled in context of an example
371
372 Fig. 1. Formal elements of Oxide and their explanation (excerpts).
373
374
375
376
377 • Provenances are divided into “concrete” ( $r$ ) and “abstract” ( $\varrho$ ). Concrete provenances are used
378 by borrow expressions, and abstract provenances are function parameters used for inputs
379 with reference type.
380
381
382 3.2 Static semantics
383 Expressions are typechecked via the judgment  $\Sigma; \Delta; \Gamma \vdash e : \tau \Rightarrow \Gamma'$ , read as: “ $e$  has type  $\tau$ 
384 under contexts  $\Sigma, \Delta, \Gamma$  producing new context  $\Gamma'$ .”  $\Delta$  contains function-level type and provenance
385 variables.  $\Gamma$  maps variables to types and provenances to pointed-to place expressions with ownership
386 qualifiers (“loans” in a “loan set” of the form  $\ell ::= \omega p$ ). For instance, when type checking  $*b := a.1$ 
387 in Figure 1b, the inputs would be  $\Delta = \bullet$  (empty) and  $\Gamma = \{a \mapsto (\text{u32}, \text{u32}), b \mapsto \&r_2 \text{ uniq u32}, r_1 \mapsto \{ \text{uniq } a.0 \}, r_2 \mapsto \{ \text{uniq } a.0 \} \}$ .
388
389 Typechecking relies on a number of auxiliary judgments, such as subtyping ( $\tau_1 \lesssim \tau_2$ ) and
390 ownership-safety ( $\Delta; \Gamma \vdash_\omega p \Rightarrow \{\bar{\ell}\}$ , read as “ $p$  has  $\omega$ -loans  $\{\bar{\ell}\}$  in the contexts  $\Delta, \Gamma$ ”). As an
391
392

```

example, consider T-ASSIGN [2019, p. 11] for the assignment expression $\pi := e$:

$$\begin{array}{c} \text{T-ASSIGN} \\ \Sigma; \Delta; \Gamma \vdash e : \tau^{\text{SI}} \Rightarrow \Gamma_1 \quad \Gamma_1(\pi) = \tau^{\text{SX}} \\ (\tau^{\text{SX}} = \tau^{\text{SD}} \vee \Delta; \Gamma_1 \vdash_{\text{uniqu}} \pi \Rightarrow \{\text{uniqu} \pi\}) \quad \Delta; \Gamma_1 \vdash \tau^{\text{SI}} \lesssim \tau^{\text{SX}} \Rightarrow \Gamma' \\ \hline \Sigma; \Delta; \Gamma \vdash \pi := e : \text{unit} \Rightarrow \Gamma'[\pi \mapsto \tau^{\text{SI}}] \triangleright \pi \end{array}$$

A valid assignment must be type-safe and ownership-safe. To be type-safe, the type of the expression τ^{SI} must be a subtype of the place's type $\Gamma_1(\pi)$. To be ownership-safe, the type must either be dead⁵, or π must have unique ownership over itself, i.e. there should be no live references to π . If so, then the type of π is updated to τ^{SI} .

3.3 Dynamic semantics

Expressions are executed via a small-step operational semantics, and the program state is a pair of a stack and an expression. A single step is represented by the judgment $\Sigma \vdash (\sigma; e) \rightarrow (\bar{\sigma}; e')$. A stack σ is a list of stack frames $\varsigma ::= \{x \mapsto v\}$ that map variables to values. For example, consider E-ASSIGN [2019, p. 16] that covers $\pi := e$ and $p := e$ expressions:⁶

$$\begin{array}{c} \text{E-ASSIGN} \\ \sigma \vdash p \Downarrow \pi^\square[x] \times \mathcal{V} \\ \hline \Sigma \vdash (\sigma; p := v) \rightarrow (\sigma[x \mapsto \mathcal{V}[v]]; ()) \end{array}$$

This rule introduces several new shorthands and administrative forms:

- The syntax $\pi^\square[x]$ means the decomposition of a place π into a root variable x and context π^\square . For example, if $\pi = a.0$ then $\pi^\square = \square.0$ and $x = a$.
- A value context \mathcal{V} is a form to handle mutation of compound objects. For instance, if $a = (0, 1)$, when evaluating $a.0 := 2$, then $\mathcal{V} = (\square, 1)$. \mathcal{V} copies all the old values, leaving a hole for the one value to be updated. Then the syntax $\mathcal{V}[v]$ means plugging v into the hole. Hence, mutating a place is represented as $\sigma[x \mapsto \mathcal{V}[v]]$.
- The judgment $\sigma \vdash p \Downarrow \pi^\square[x] \times \mathcal{V}$ evaluates a place expression under the current stack into a place π and value context \mathcal{V} . For instance, if p is a dereference of a reference, then this judgment resolves p to the concrete memory location π it points-to under σ .

3.4 Formalized principles

Now, we have enough of the language formalized to give a precise statement of each slicing principle from Section 2. Each principle will be presented with the corresponding theorem, using underlining in color to highlight correspondences.

In the principles and corresponding algorithm/proofs, there are many concepts which we distinguish by notational convention. We denote objects by their metavariable, e.g. p or σ , and add a sans-serif subscript for distinct roles where needed, e.g. π_{mut} for a mutated place and π_{any} for an arbitrary place. We generally use a superscript i for an object that varies between two executions

⁵Oxide uses the metavariables τ^{SD} to mean “dead types” and τ^{SX} to mean “possibly dead types”. A place becomes dead when it is moved, e.g. see T-MOVE in [2019, p. 11]. T-ASSIGN allows a dead place to be revived. For instance, consider the program:

```
let x : String = "a"; print(x); x := "b"
```

When $\text{print}(x)$ moves x , its type is updated to String^\dagger in Γ . Then the T-ASSIGN rule permits x to be assigned again to “revive” that place, setting its type back to String .

⁶This E-ASSIGN rule is not the exact same rule that appears in Weiss et al. [2019, p. 16], as the published version is incorrect. In correspondence with the authors, we determined that the rule presented here has the intended semantics. Additionally, we do not use the referent \mathcal{R} construct of Oxide since we do not consider arrays in this paper, so we use π anywhere \mathcal{R} would otherwise appear.

of a program, like σ^i or v^i . And we use right arrows to indicate changes to an object after stepping (instead of primes, to avoid polluting the superscript), e.g. σ^i versus $\vec{\sigma}^i$.

444

445

PRINCIPLE 1 (SLICING PRINCIPLE FOR PLACES)

A mutation to a place is a mutation to all conflicting places.

THEOREM 3.1. Let:

- $\pi_{\text{mut}} = \pi_{\text{mut}}^\square[x]$, σ where $\sigma \vdash \pi_{\text{mut}}^\square \times x \Downarrow \mathcal{V}$
- $v, \vec{\sigma} = \sigma[x \mapsto \mathcal{V}[v]]$
- π_{any} be any place

Then $\sigma(\pi_{\text{any}}) \neq \vec{\sigma}(\pi_{\text{any}}) \implies \pi_{\text{any}} \sqcap \pi_{\text{mut}}$

451

452

As described in Section 3.3, a mutation to a place is represented by updating a variable x in a stack σ by plugging a value v into a value context \mathcal{V} . To denote a conflict, we reuse the notation from Oxide that $\pi_1 \# \pi_2$ means “ π_1 and π_2 do not conflict”, or more formally:

453

454

455

$$x_1.q_1 \# x_2.q_2 \stackrel{\text{def}}{=} x_1 \neq x_2 \vee ((q_1 \text{ is not a prefix of } q_2) \wedge (q_2 \text{ is not a prefix of } q_1))$$

456

457

458

Conversely, we use the shorthand $\pi_1 \sqcap \pi_2 \stackrel{\text{def}}{=} \neg(\pi_1 \# \pi_2)$. So if a place π_{any} is changed when π_{mut} is mutated, then it must be that $\pi_{\text{any}} \sqcap \pi_{\text{mut}}$.

459

460

461

462

463

PRINCIPLE 2 (SLICING PRINCIPLE FOR REFERENCES)

The lifetime of a reference contains all potential aliases of what the reference points-to.

THEOREM 3.2. Let:

- σ where $\Sigma \vdash \sigma : \Gamma$
- p_{mut} where $\bullet; \Gamma \vdash_{\text{uniq}} p_{\text{mut}} \Rightarrow \{\bar{\ell}\}$ and $\sigma \vdash p_{\text{mut}} \Downarrow \pi_{\text{mut}}$
- p_{any} where $\sigma \vdash p_{\text{any}} \Downarrow \pi_{\text{any}}$

Then $\pi_{\text{any}} \sqcap \pi_{\text{mut}} \implies \exists^{\text{uniq}} p_{\text{loan}} \in \{\bar{\ell}\}. p_{\text{any}} \sqcap p_{\text{loan}}$

464

465

466

Rather than referring to a lifetime directly, we instead use Oxide’s ownership safety judgment described in Section 3.2 to get the corresponding loan set for a mutated place expression p_{mut} . If p_{mut} includes a dereference, then the loan set should include potential aliases.

467

468

469

A notable detail is that we do not compare the loan sets of p_{mut} and p_{any} to see if they contain conflicting places, but rather compare p_{any} just against the loan set of p_{mut} . This works because the loan set contains not just the set of places π that p_{mut} could point-to, but also the set of other references to the places p_{mut} points-to (via reborrows).

470

471

472

473

474

475

476

477

478

479

480

481

482

483

484

485

486

487

PRINCIPLE 3 (SLICING PRINCIPLE FOR FUNCTION CALLS))

When calling a function:

(a) only mutable references in the arguments can be mutated, and ...

THEOREM 3.A. Let:

- $\Gamma, \pi_{\text{arg}}, \sigma$ where $\Gamma(\pi_{\text{arg}}) = \tau^{SI}$ and $\Sigma \vdash \sigma : \Gamma$
- f where $\Sigma \vdash (\sigma; f(\pi_{\text{arg}})) \xrightarrow{*} (\vec{\sigma}; _)$
- $\vec{\sigma}' = \vec{\sigma}[\forall p_{\text{ref}} \in \text{uniq-refs}(\pi_{\text{arg}}, \tau^{SI}) . p_{\text{ref}} \mapsto \sigma(p_{\text{ref}})]$

Then $\sigma = \vec{\sigma}'$.

488

489

First, we define “mutable references in the arguments” as $\omega\text{-refs}(p, \tau)$ that returns the ω -safe place expressions of references inside of p of type τ . For instance, if $x = 0$ and $y = (0, \&r \text{ uniq } x)$

490

491 then $\text{uniq-refs}(y, (\text{u32}, \&r \text{ uniq u32})) = \{*\text{y}.\text{1}\}$. The definition is:

$$\omega\text{-refs}(p, \tau^B) = \emptyset \quad \omega\text{-refs}(p, (\tau_1^{SI}, \dots, \tau_n^{SI})) = \bigcup_i \omega\text{-refs}(p.i, \tau_i^{SI})$$

$$\omega\text{-refs}(p, \&\rho \omega' \tau^{XI}) = \begin{cases} \{*\text{p}\} \cup \omega\text{-refs}(*\text{p}, \tau^{XI}) & \text{if } \omega' \lesssim \omega^{\text{(7)}} \\ \emptyset & \text{otherwise} \end{cases}$$

497 Then we define Theorem 3.A in the theme of a transaction: if all the changes to unique references
498 in π_{arg} are rolled back, then the new stack is the same as the one before the function call. This
499 means implicitly that no other values could have been mutated.

500

502 PRINCIPLE 3 (SLICING PRINCIPLE FOR
503 FUNCTION CALLS)

504 When calling a function:

505 (b) ...the mutations and return value
506 are only influenced by the arguments.

507 THEOREM 3.B. Let:

- $\Gamma, \pi_{\text{arg}}, \sigma^i$ where $\Gamma(\pi_{\text{arg}}) = \tau^{SI}$ and
 $i \in \{1, 2\}$ and $\Sigma \vdash \sigma^i : \Gamma$
 - f where $\Sigma \vdash (\sigma^i; f(\pi_{\text{arg}})) \xrightarrow{*} (\vec{\sigma}^i; v^i)$
 - $P = \text{all-places}(\pi_{\text{arg}}, \tau^{SI})$
- 508 Then $\sigma^1 \sim_P \sigma^2 \implies \vec{\sigma}^1 \sim_P \vec{\sigma}^2 \wedge v^1 = v^2$

510 The idea behind Theorem 3.B is that “influence” is translated into a form of noninterference:
511 if the input to a function is the same under any two stacks σ^1 and σ^2 , then the mutations to that
512 input must be the same. The rest of the stack is allowed to vary, but because the function f cannot
513 read it, that variation cannot influence the final value.

514 To formalize “the input being the same”, we introduce another auxiliary function for transitive
515 equality. For instance, if we only required that $\sigma^1(\pi_{\text{arg}}) = \sigma^2(\pi_{\text{arg}})$ where $\pi_{\text{arg}} = \text{ptr } x$, then if
516 $\sigma^1(x) \neq \sigma^2(x)$ the theorem would not be true. Hence, transitive equality is defined as equality
517 including all pointed places. We define this concept through two pieces: a function for generating
518 the set of places (denoted by P), and a relation defining the equivalence of stacks for a set of places.

$$\text{all-places}(p, \tau^{SI}) \stackrel{\text{def}}{=} \{p\} \cup \text{shrd-refs}(p, \tau^{SI})$$

$$\sigma^1 \sim_P \sigma^2 \stackrel{\text{def}}{=} \forall p \in P . \sigma^1(p) = \sigma^2(p)$$

523 Therefore Theorem 3.B states that if π_{arg} is transitively equal under two otherwise arbitrary stacks,
524 then π_{arg} is still transitively equal after evaluating $f(\pi_{\text{arg}})$, and the output of $f(\pi_{\text{arg}})$ is also equal.

525 4 ALGORITHM

527 Now, we use the formalized principles to build an algorithm for slicing Oxide. A slicing algorithm
528 canonically has the form of a function that takes as input a program and a slicing criterion, and
529 outputs a subset of the program [Weiser 1984]. However, Abadi et al. [1999] observed that a slice is a
530 particular perspective on a term’s dependencies, and so an alternative view on slicing is as a system
531 for tracking dependencies. The slicing calculus of Abadi et al. has a type judgment of the form
532 $\Gamma \vdash e : (\tau, \kappa)$ where expressions have a type τ and set of dependencies κ . Then, type-checking a
533 program involves computing the dependencies of all terms, and computing a slice is simply looking
534 up the dependencies in the term’s type.

535 Inspired by Abadi et al., we formulate our algorithm as an extension to the Oxide type system
536 that calculates dependencies for all expressions and places in a program. First, we assume that all

537 ⁷The relation $\omega' \lesssim \omega$ (again re-used from Oxide) is defined as $\text{uniq} \not\leq \text{shrd}$ and otherwise $\omega' \lesssim \omega$. This relation captures
538 the concept that a unique reference can be used as a shared reference, but not the other way around.

540	$\boxed{\Sigma; \Delta; \Gamma; \Theta \vdash e : \tau \bullet \kappa \Rightarrow \Gamma'; \Theta'}$
541	T-u32
542	
543	$\Sigma; \Delta; \Gamma; \Theta \vdash n_l : \text{u32} \bullet \{l\} \Rightarrow \Gamma; \Theta$
544	
545	T-LET
546	$\Sigma; \Delta; \Gamma; \Theta \vdash e_1 : \tau_1^{\text{SI}} \bullet \kappa_1 \Rightarrow \Gamma_1; \Theta_1 \quad \Gamma; \Delta_1 \vdash \tau_1^{\text{SI}} \lesssim \tau_a^{\text{SI}} \Rightarrow \Gamma'_1$
547	$\Theta'_1 = \Theta_1[\forall \pi^\square[x] . \pi^\square[x] \mapsto \kappa_1] \quad \Sigma; \Delta; \text{gc-loans}(\Gamma'_1, x : \tau_a^{\text{SI}}); \Theta'_1 \vdash e_2 : \tau_2^{\text{SI}} \bullet \kappa_2 \Rightarrow \Gamma_2, x : \tau^{\text{SP}}; \Theta_2$
548	<hr/>
549	$\Sigma; \Delta; \Gamma; \Theta \vdash \text{let } x : \tau_a^{\text{SI}} = e_1; e_2 : \tau_2^{\text{SI}} \bullet \kappa_2 \Rightarrow \Gamma_2; \Theta_2$
550	
551	T-MOVE
552	$\Delta; \Gamma \vdash_{\text{uniq}} \pi \Rightarrow \{ \text{uniq} \pi \} \quad \Gamma(\pi) = \tau^{\text{SI}} \quad \text{noncopyable}_{\Sigma} \tau^{\text{SI}}$
553	$\Sigma; \Delta; \Gamma; \Theta \vdash \pi : \tau^{\text{SI}} \bullet \Theta(\pi) \Rightarrow \Gamma[\pi \mapsto \tau^{\text{SI}^\dagger}]; \Theta$
554	
555	T-ASSIGN
556	$\Sigma; \Delta; \Gamma; \Theta \vdash e : \tau^{\text{SI}} \bullet \kappa_e \Rightarrow \Gamma_1; \Theta_1 \quad \Gamma_1(\pi) = \tau^{\text{SX}} \quad (\tau^{\text{SX}} = \tau^{\text{SD}} \vee \Delta; \Gamma_1 \vdash_{\text{uniq}} \pi \Rightarrow \{ \text{uniq} \pi \})$
557	$\Delta; \Gamma_1 \vdash \tau^{\text{SI}} \lesssim \tau^{\text{SX}} \Rightarrow \Gamma' \quad \Theta'_1 = \Theta_1[\text{update-conflicts } (\Theta_1, \pi, \kappa_e)]$
558	<hr/> $\Sigma; \Delta; \Gamma; \Theta \vdash \pi := e : \text{unit} \bullet \emptyset \Rightarrow \Gamma'[\pi \mapsto \tau^{\text{SI}}] \triangleright \pi; \Theta'_1$
559	
560	T-ASSIGNDEREF
561	$\Sigma; \Delta; \Gamma; \Theta \vdash e : \tau_n^{\text{SI}} \bullet \kappa_e \Rightarrow \Gamma_1; \Theta_1 \quad \Delta; \Gamma_1 \vdash_{\text{uniq}} p : \tau_o^{\text{SI}} \quad \Delta; \Gamma_1 \vdash_{\text{uniq}} p \Rightarrow \{\bar{\ell}\}$
562	$\Delta; \Gamma_1 \vdash \tau_n^{\text{SI}} \lesssim \tau_o^{\text{SI}} \Rightarrow \Gamma' \quad \Theta'_1 = \Theta_1[\forall \omega p' \in \{\bar{\ell}\} . \text{update-conflicts } (\Theta_1, p', \kappa_e)]$
563	<hr/> $\Sigma; \Delta; \Gamma; \Theta \vdash p := e : \text{unit} \bullet \emptyset \Rightarrow \Gamma' \triangleright p; \Theta'_1$
564	
565	T-APP
566	$\overline{\Sigma; \Delta; \Gamma \vdash \Phi} \quad \overline{\Delta; \Gamma \vdash \rho} \quad \overline{\Sigma; \Delta; \Gamma \vdash \tau^{\text{SI}}}$
567	$\Sigma(f) = \text{fn } f \langle \bar{\psi}, \bar{\varrho}, \bar{\alpha} \rangle(x : \tau_a^{\text{SI}}) \rightarrow \tau_r^{\text{SI}} \text{ where } \overline{\varrho_1 : \varrho_2} \{ e \}$
568	$\Sigma; \Delta; \Gamma; \Theta \vdash \pi : \tau_a^{\text{SI}}[\Phi/\varrho][\rho/\varrho][\tau^{\text{SI}}/\alpha] \bullet \kappa \Rightarrow \Gamma_1; \Theta_1 \quad \Delta; \Gamma_1 \vdash \varrho_2[\rho/\varrho] \triangleright \varrho_1[\rho/\varrho] \Rightarrow \Gamma_2$
569	$\Theta_2 = \Theta_1[\forall p \in \text{uniq-refs}(\pi, \tau_a^{\text{SI}}) . \Delta; \Gamma_2 \vdash_{\text{uniq}} p \Rightarrow \{\bar{\ell}\}, \forall \omega p' \in \{\bar{\ell}\} . \text{update-conflicts } (\Theta_1, p', \kappa)]$
570	<hr/> $\Sigma; \Delta; \Gamma; \Theta \vdash f \langle \bar{\Phi}, \bar{\varrho}, \bar{\tau^{\text{SI}}} \rangle(\pi) : \tau_r^{\text{SI}}[\Phi/\varrho][\rho/\varrho][\tau^{\text{SI}}/\alpha] \bullet \kappa \Rightarrow \Gamma_2; \Theta_2$
571	

Fig. 2. Modular slicing algorithm formulated as an extension to Oxide’s static semantics. Dependencies for expressions are represented as $\bullet \kappa$, and dependencies for places are contained in Θ . Remaining rules are given in Appendix A.1.

Oxide expressions e are uniquely labeled with their location l (written as e_l). Second, we modify the type-checking judgment to include two new features: a dependency set $\kappa = \{\bar{l}\}$ for the value of an expression, and a dependency context $\Theta = \{\pi \mapsto \kappa\}$ for the dependencies of places in memory. As a convention, we will highlight in red our additions to the rules, and leave in black the original logic from Weiss et al. [2019]. So the new judgment has the form:

$$\Sigma; \Delta; \Gamma; \Theta \vdash e : \tau \bullet \kappa \Rightarrow \Gamma'; \Theta'$$

This judgment reads as: “under global context Σ , function context Δ , type context Γ , and dependency context Θ , expression e has type τ and dependencies κ and produces a new type context Γ' and dependency context Θ' .”

589 4.1 Rules

590 Figure 2 introduces a subset of the modified static semantics for dependency tracking. The general
 591 goal is that for an expression e , κ contains the set of dependencies needed to determine the value
 592 of just e . So for instance, if e has type unit, then $\kappa = \emptyset$ because it has no distinguishable values.
 593 Then for Θ we have a corresponding entry for every place π in Γ . As places are mutated, additional
 594 dependencies are added to $\Theta(\pi)$ based on the dependencies of the mutation's inputs.

595 The base case is that constructors initialize the slice set with their location. For example, with
 596 rule T-U32, a labeled number n_l has the single dependency $\{l\}$, and the dependency context Θ is
 597 unchanged. Appendix A.1 shows additional constructors like T-TUPLE and T-BORROW.

598 Places are introduced in T-LET with two notable aspects: first, dependency contexts are threaded
 599 through expressions in the order of evaluation. So e_1 produces Θ_1 which is then passed to checking
 600 of e_2 . Second, x is bound to e_1 with dependencies κ_1 , so that fact is saved in the updated dependency
 601 context Θ'_1 as $x \mapsto \kappa_1$. For compound objects, an entry in the dependency context is created for
 602 every path into x , represented as $\pi^\square[x]$. Then a read of a place π has the dependencies $\Theta(\pi)$ as in
 603 T-MOVE

604 T-ASSIGN handles assignments to places π , i.e. without dereferences. The overall expression has
 605 unit value and hence has an empty dependency set. The main step is to update Θ to reflect the
 606 dependencies κ_e of e . The naive implementation $\Theta[\pi \mapsto \kappa_e]$ doesn't work because of Principle 1:
 607 all conflicting places are also mutated, where conflicts are formally represented as $p_1 \sqcap p_2$. Hence,
 608 we define a function update-conflicts that updates all the conflicts of a place expression p with κ
 609 in Θ :

$$610 \text{update-conflicts}(\Theta, p, \kappa) \stackrel{\text{def}}{=} \forall \pi \mapsto \kappa_\pi \in \Theta . (p \sqcap \pi \implies \pi \mapsto \kappa_\pi \cup \kappa)$$

612 T-ASSIGNDeref generalizes to assignments with dereferences. As per Principle 2, any place pointed
 613 by p is contained in the loan set $\{\bar{l}\}$. Therefore this rule propagates the dependencies κ_e of e to
 614 every such place π in the dependency context Θ . Additionally, $\{\bar{l}\}$ contains reborrowed pointers of
 615 the form $*\pi$ which also get their dependencies updated.

616 Finally, function calls are handled by T-APP.⁸ Each place in the loan set of each unique reference in
 617 uniq-refs(π, τ^{SI}) is updated with the dependencies of the argument κ . For example, in the program:

618 `let x : u32 = 0; let y : &r uniq u32 = &r uniq x; let z : (u32, &r uniq u32) = (1l, y); f(z)`

620 The location l of the number 1 in z should be included in the dependencies of x , since f could
 621 use $z.0$ in mutating $*z.1$. This happens because $\text{uniq-refs}(z, (u32, &r \text{uniq } u32)) = \{ *z.1 \}$, and
 622 $\bullet ; \Gamma \vdash_{\text{uniqu}} *z.1 \Rightarrow \{x\}$. Then because $\kappa = \Theta(z)$ and $l \in \Theta(z)$, the T-App rule ensures that l is added
 623 to $\Theta(x)$.

624 4.2 Soundness

626 To characterize the correctness of our algorithm, we again draw on the insight of Abadi et al. [1999]:
 627 our algorithm is sound if it satisfies a form of noninterference. At a high-level, if an expression e
 628 has some dependencies κ , then for two stacks σ^1, σ^2 , if those stacks are the same for at least κ , then
 629 e should evaluate to the same value whether it is run with σ^1 or σ^2 . To define "same for at least κ ",
 630 we first map from locations in κ to places in Θ as follows:

$$631 \text{deps}(\Theta, \kappa) \stackrel{\text{def}}{=} \{ \pi \mid \pi \mapsto \kappa_\pi \in \Theta \wedge \kappa_\pi \subseteq \kappa \}$$

633 ⁸Besides E-ASSIGN, T-APP is the only other rule that is significantly modified from the published version of Oxide. Originally,
 634 T-APP permitted a list of argument expressions. For the sake of simplifying the slicing algorithm and proof, we restricted
 635 functions to only take a single argument which must be a place, not an expression. We also restricted functions to only be
 636 top-level, and not closures.

Then if $P = \text{deps}(\Theta, \kappa)$, two stacks are equivalent under P if $\sigma^1 \sim_P \sigma^2$. We will use the shorthand $\sigma^1 \sim_{\kappa}^{\Theta} \sigma^2 \stackrel{\text{def}}{=} \sigma^1 \sim_{\text{deps}(\Theta, \kappa)} \sigma^2$.

Because our language has effects, noninterference must also account for them: for all places π in the stack, if the stacks are equivalent for the place's dependencies κ_{π} , then π should be the same in the output stacks after executing e . We formalize these statements as a noninterference theorem:

THEOREM 4.1 (NONINTERFERENCE). *Let e such that $\Sigma; \bullet; \Gamma; \Theta \vdash e : \tau \bullet \kappa \Rightarrow \Gamma'; \Theta'$. For $i \in \{1, 2\}$, let σ_i such that $\Sigma \vdash \sigma_i : \Gamma$ and $\Sigma \vdash (\sigma_i; e) \xrightarrow{*} (\vec{\sigma}^i; v^i)$. Then:*

- (a) $\sigma^1 \sim_{\kappa}^{\Theta} \sigma^2 \implies v^1 = v^2$
- (b) $\forall \pi \mapsto \kappa_{\pi} \in \Theta'. \sigma^1 \sim_{\kappa_{\pi}}^{\Theta} \sigma^2 \implies \vec{\sigma}^1(\pi) = \vec{\sigma}^2(\pi)$

Note that the context Δ is required to be empty because an expression e can only evaluate if it does not contain abstract type or provenance variables. The judgment $\Sigma \vdash \sigma^i : \Gamma$ means “the stack σ^i is well-typed under Σ and Γ ”. That is, all places π in Γ must exist and have values of the corresponding type in σ .

Additionally, we ignore the issue of non-termination in this definition. Distinguishing terminating from non-terminating programs is important in the context of security, but less so for slicing. So we make the simplifying assumption that the theorem only applies to terminating programs.

PROOF SKETCH. The proof proceeds by induction on the derivation of the typing judgment. The full proof is in Appendix A.2, and here we will sketch the proof of part (b) on the key cases. Most of the novel pieces relate to effects on the stack, so the proof of (a) is not as interesting.

Let $\pi_{\text{any}} \mapsto \kappa_{\text{any}} \in \Theta'$. Assume the initial heaps from the two runs are equivalent under κ_{any} , i.e. $\sigma^1 \sim_{\kappa_{\text{any}}}^{\Theta} \sigma^2$. Then we want to show $\vec{\sigma}^1(\pi_{\text{any}}) = \vec{\sigma}^2(\pi_{\text{any}})$.

T-ASSIGN: $e = "\pi_{\text{mut}} := e"$. First, $\sigma^1 \sim_{\kappa_{\text{any}}}^{\Theta} \sigma^2$ implies π_{any} is initially equal between the two runs, i.e. $\sigma^1(\pi_{\text{any}}) = \sigma^2(\pi_{\text{any}})$. Then by the inductive hypothesis on e , π_{any} is still equal after stepping in both runs (σ^1, e) to $(\vec{\sigma}_e^1, v_e^1)$ and (σ^2, e) to $(\vec{\sigma}_e^2, v_e^2)$. Finally, the runs conclude by updating the place π_{mut} to v_e^i . If π_{mut} has root variable x , i.e. $\pi_{\text{mut}} = \pi^{\square}[x]$, then updating π_{mut} is done by computing the value of x without π_{mut} as \mathcal{V}^i : $\vec{\sigma}_e^i \vdash \pi_{\text{mut}} \Downarrow \mathcal{V}^i$. Then the stack is updated by plugging v_e^i into the context: $\vec{\sigma}^i = \vec{\sigma}_e^i[x \mapsto \mathcal{V}^i[v_e^i]]$.

By Theorem 3.1, the only way π_{any} could have changed is if π_{mut} and π_{any} conflict, i.e. $\pi_{\text{mut}} \sqcap \pi_{\text{any}}$. By the premises of T-ASSIGN, then $\Theta'(\pi_{\text{any}}) = \Theta(\pi_{\text{any}}) \cup \kappa_e$ where κ_e are the dependencies of e . Because $\kappa_e \subseteq \Theta'(\pi_{\text{any}})$ then $\sigma^1 \sim_{\kappa_e}^{\Theta} \sigma^2$, which by the inductive hypothesis (a) on e means $v_e^1 = v_e^2$, and subsequently $\vec{\sigma}^1(\pi_{\text{any}}) = \vec{\sigma}^2(\pi_{\text{any}})$.

T-ASSIGNDeref: $e = "p_{\text{mut}} := e"$. This proof is similar to T-ASSIGN, except we must account for whether p_{mut} could point to a conflicting place with π_{any} . That is, where $\vec{\sigma}_e^i \vdash p_{\text{mut}} \Downarrow \pi_{\text{mut}}^i$ and $\pi_{\text{any}} \sqcap \pi_{\text{mut}}^i$. By Theorem 3.2, because π_{mut}^i and π_{any} conflict, then π_{any} must conflict with a place in the loan set of p_{mut} . Therefore by T-ASSIGNDeref, $\kappa_e \subseteq \kappa_{\text{any}}$, and $\vec{\sigma}^1(\pi_{\text{any}}) = \vec{\sigma}^2(\pi_{\text{any}})$.

T-APP: $e = "f(\pi_{\text{arg}})"$. Then $(\sigma^i, f(\pi_{\text{arg}})) \xrightarrow{*} (\vec{\sigma}^i, v^i)$. By Theorem 3.A, if $\sigma^i(\pi_{\text{any}}) \neq \vec{\sigma}^i(\pi_{\text{any}})$, then π_{any} conflicts with a unique reference p_{ref} projected from π_{arg} . Then by Theorem 3.2, π_{any} is in the loan set of p_{arg} , and by T-APP then π_{arg} is registered as a dependency of π_{any} , i.e. $\kappa_{\text{arg}} \subseteq \kappa_{\text{any}}$. Because σ^1 and σ^2 are then equivalent up to κ_{arg} , then π_{arg} is initially equal between stacks. By Theorem 3.B, π_{arg} is also equal after executing f . Therefore no mutations could have occurred, and $\vec{\sigma}^1(\pi_{\text{any}}) = \vec{\sigma}^2(\pi_{\text{any}})$ \square

687 5 IMPLEMENTATION

688 We have established that our modular slicing algorithm is sound, but it is still an approximation to
 689 the true slice. For example, a function could take a unique reference as input, but then not mutate
 690 it. Hence, our next goal is to evaluate the *precision* of our algorithm: how much larger are its slices
 691 than optimal?

692 We chose to address this question empirically, by generating slices on real-world codebases and
 693 comparing the slice sizes. To some extent, the practical precision of the slicer depends upon the
 694 frequency with which programmers write code where the slicer's modular assumptions are too
 695 strong. Therefore we needed an implementation of the slicing algorithm that works on real code.
 696 Using Oxide to verify the correctness of our algorithm, we then translated the key concepts into a
 697 slicer that works for Rust.

698 Rust imposes a number of logistical constraints that required us to adapt the ideas above such
 699 that they could be integrated with the Rust compiler. Below, we discuss the implementation details
 700 needed to meet these constraints. First, Rust computes lifetime information not on the surface
 701 language like Oxide, but on an LLVM-like intermediate representation called MIR (5.1). Second,
 702 Rust does not use the same loan-set model of lifetimes as Oxide, but rather a line-set model (5.2).
 703 Finally we describe a few details necessary to execute the slicer on Rust programs (5.3).

705 5.1 MIR

706 The Rust compiler lowers programs into a “mid-level representation”, or MIR, that represents
 707 programs as a control-flow graph. Essentially, expressions are flattened into sequences of instruc-
 708 tions (basic blocks) which terminate in instructions that can jump to other blocks, like a branch or
 709 function call.

710 To implement the modular slicing algorithm for MIR, we largely reused standard techniques
 711 for analyzing CFGs. Specifically, we implemented a dataflow analysis that takes as input a slicing
 712 criterion, and finds all relevant instructions in the CFG to the criterion by maintaining a set of
 713 relevant places. The analysis is:

- 715 • *Backwards*, e.g. if p is determined to be relevant after $p := e$, then e is registered as relevant
 716 before $p := e$.
- 717 • *Flow-sensitive*, i.e. a different set of relevant places is maintained at each instruction.
- 718 • *Gen-kill*, i.e. each instruction is processed by a transfer function that either adds or removes
 719 places in the relevant place set.

720 The analysis is iterated to a fixpoint over the set of relevant places. The transfer function largely
 721 implements the same transformations as the algorithm in Section 4. For instance, the transfer
 722 function for assignment instructions, shown here in pseudocode, looks similar to T-ASSIGNDEREF.

```
724 1 let instruction = (p := e);
 2 if ∃ p1 ∈ loans(p), p2 ∈ relevant set . p1 ⊑ p2 {
 3   relevant set ∪= places(e)
 4 }
```

To handle conditional effects as described in Section 2.1, we compute the control-dependence between instructions. If a relevant instruction is control-dependent

729 on a branch, then the branch is also relevant. We define control-dependence following Ferrante
 730 et al. [1987]: an instruction I_X is control-dependent on I_Y if there exists a directed path P from
 731 I_X to I_Y such that any I_Z in P is post-dominated by I_Y , and I_X is not post-dominated by I_Y . An
 732 instruction I_X is post-dominated by I_Y if I_Y is on every path from I_X to a return node. We compute
 733 control-dependencies by generating the post-dominator tree and frontier of the CFG using the
 734 algorithms of Cooper et al. [2001] and Cytron et al. [1989], respectively.

```

736
737
738
739 1 fn get_count(
740 2     h: &mut HashMap<String, u32>,
741 3     k: String
742 4 ) -> u32 {
743 5     if !h.contains_key(&k) {
744 6         h.insert(k, 0);
745 7         0
746 8     } else {
747 9         *h.get(&k).unwrap()
748 10    }
749 11 }
750
751
752
753

```

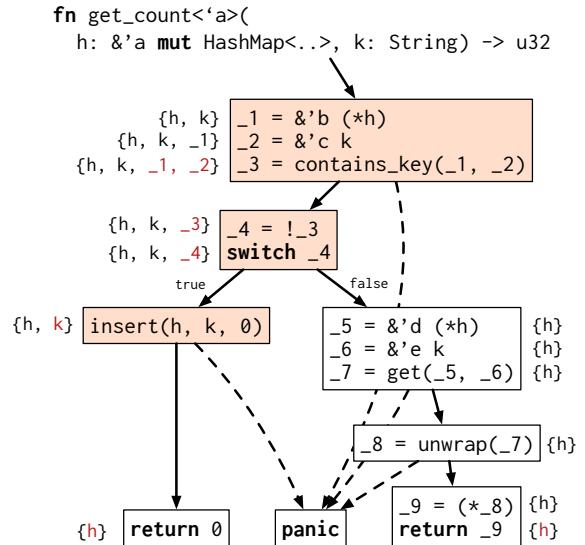


Fig. 3. Example of a Rust function (left) and corresponding MIR graph (right) sliced on the variable h at the end of the function. The MIR graph shows the result of the slicing dataflow analysis after reaching a fixpoint. Relevant MIR instructions are highlighted in orange, and the relevant place sets are adjacent to instructions with additions highlighted in red.

Figure 3 shows an example of the Rust slicer in action. First, it demonstrates how the function `get_count` is lowered from Rust to MIR. For instance, the condition on `contains_key` is lowered to a `switch` statement. Second, it shows how a relevant place set is maintained at each node, and that instructions which modify the place set are considered relevant.

5.2 Lifetimes

Because Rust does not compute loans for lifetimes, we had to regenerate the loans ourselves. The key observation is that both Oxide's borrow checker and Rust's borrow checker share a common ingredient: subtyping in the form of outlives-constraints. An outlives-constraint is a relation between lifetimes. Constraints are generally created either implicitly via subtyping, or explicitly via user-provided constraints in a function's signature. For instance, consider an assignment like `let x: &'a i32 = &'b y`. In both languages, type-checking this assignment requires $\text{'b i32} \leq \text{'a}$ which in turn requires 'b : 'a .

In Oxide, the act of proving subtyping between lifetimes causes the loan sets to change. Specifically, Rule OL-LOCALPROVENANCES [2019, p. 14] states that proving $r_1 :> r_2$ updates $\Gamma(r_2) = \Gamma(r_1) \cup \Gamma(r_2)$. From this rule, we can generalize to the principle: if $\text{loans('b)} \subseteq \text{loans('a)}$ then 'b : 'a . Hence, we modified the Rust compiler to export the set of outlives-constraints generated by a program, and computed the reflexive-transitive closure to generate a complete list of outlives-constraints. Then, we compute the loans for a lifetime as follows:

```

780
781 1 loans('a) = {p} if expression "&'a p" exists in the program
782 2 loans('a) =  $\bigcup_{\text{'b : 'a}} \text{loans('b)}$ 
783
784

```

```

785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833

```

```

fn main() {
    let mut buffer: String = String::new();
    let stdin: Stdin = io::stdin();

    let mut count: i32 = 0;
    loop {
        buffer.clear();
        stdin.read_line(&mut buffer).unwrap();
        let input: &str = buffer.trim();

        count += 1;
        println!("read {} on iteration {}", input, count);
    }
}

```

Fig. 4. Screenshot of our VS Code extension that uses the modular slicer to highlight relevant lines of code for a given criterion. The variable `input` is sliced on line 14.

5.3 Running the slicer

To slice a given Rust program, we take as input a slicing criterion, then run the Rust compiler to generate MIR and lifetime constraints for the function containing the criterion. We then run the slicing dataflow analysis seeded by the criterion.

For the evaluation, we simply count the number of relevant instructions for a given slice. For usage of the tool beyond the context of the paper, we have also developed an IDE extension for VSCode that source-maps relevant instructions back to spans of code in the surface program. Figure 4 shows an example program sliced within the GUI.

Our fork of the Rust compiler (to export lifetime outlives-constraints) and our modular slicer are both free and open-source. A GitHub link will be provided upon publication, and for double-blind review, an anonymized copy of the source code is provided in the supplementary materials.

6 EVALUATION

To evaluate the precision of modular slicing, we ran the modular Rust slicer over a dataset of Rust codebases. We then compared the size of slices for a given criterion against a series of alternatives. Specifically, we evaluate in two directions:

- (1) What if the algorithm had *more* information? If the slicer had access to the definition of called functions, how much more precise are the whole-program slices vs. an equivalent modular slice?
- (2) What if the algorithm had *less* information? If Rust's type system was more like C, i.e. it lacked ownership, then how much less precise do the modular slices become?

To answer these questions, we modified the slicer described in Section 5 in three ways:

- **WHOLE-PROGRAM:** the slicer will recursively slice function calls. For example, if calling a function `f(&mut x)` where `f` does not actually modify `x`, then the recursive slicer will recognize `x` is not mutated while the modular slicer will assume that `x` is mutated.

- MUT-BLIND: the slicer does not distinguish between mutable and immutable references. For example, if calling a function `f(&x)`, then the slicer assumes that `x` can be modified.
- REFERENCE-BLIND: the slicer does not use lifetimes to reason about references, and rather assumes all references of the same type can alias. For example, if a function takes as input `f(x: &mut i32, y: &mut i32)` then `x` and `y` are assumed to be aliases.

The MUT-BLIND and REFERENCE-BLIND modifications represent an ablation of the precision provided by ownership types. The WHOLE-PROGRAM modification represents the most precise slicer we can feasibly implement. Each modification can be combined with the others, representing $2^3 = 8$ possible conditions for evaluation.

WHOLE-PROGRAM requires some clarification on its implementation: it works by transferring slicing criteria from the call site to the function definition. For example, if a slice in function `f` calls function `g(&mut y)` where `y` is relevant, then the slicer checks whether `g` ever modifies its first input, and consider the call relevant if so.

Due to the architecture of the Rust compiler, it is difficult to analyze more than one package, or “crate”, at a time. Hence WHOLE-PROGRAM can only recurse *within a given package*, and not into the standard library or other dependencies. If WHOLE-PROGRAM reaches a function call that cannot be recursed, then it falls back onto the modular slicer. This also includes situations like higher-order functions where no definition is available.

With these three modifications, we compare the size of slices (6.1) on a dataset of Rust projects (6.2) to quantitatively (6.3) and qualitatively (6.4) evaluate the precision of the modular slicer.

6.1 Metric

To compare the slices generated in different conditions, we need a metric for precision. As an example, consider this slice on a vector:

<pre> 559 // Modular slice 560 let mut v = vec![1, 2]; 561 let x = v.get_mut(0).unwrap(); 562 println!("{} {}", v, *x); </pre>	<pre> 1 // Whole-program slice 2 let mut v = vec![1, 2]; 3 let x = v.get_mut(0).unwrap(); 4 println!("{} {}", v, *x); </pre>
--	--

The function `Vec::get_mut` returns a mutable reference to an element of the vector. Because it takes as input `&mut self`, the modular slicer assumes that `get_mut` mutates `v`, even though it is simply passing the mutable permission to an inner element. Therefore the modular slice (left) is larger than the whole-program slice (right).

To measure the difference between these slices, we need to compute and compare their sizes. The modular slice has 2 lines and the true slice has 1 line. To control for the size of the program being sliced, we use the percentage difference between the slice sizes: here, the modular slice is 100% larger. For Rust slices generated on the MIR, we compare the number of instructions within each slice.

6.2 Dataset

To create a dataset of slices, we first needed a dataset of Rust crates to slice. We had two main selection criteria: first, due to the single-crate limitation of WHOLE-PROGRAM, we preferred large crates so as to see a greater impact from WHOLE-PROGRAM, specifically crates with over 10,000 lines of code as measured by the `cloc` utility [Danial 2021]. Second, to control for code styles specific to individual applications, we wanted crates from a wide range of domains.

After a manual review of large crates in the Rust ecosystem, we selected 10 crates, shown in Table 1. We built each crate with as many feature flags enabled as would work on our Ubuntu 16.04 machine. Additional details like the specific flags and exact commit hashes used can be found in

Project	Crates	Purpose	LOC	# Slices	Avg. Slices/Func
rayon		Data parallelism library	15,524	10,535	9.9
Rocket	core/lib	Web backend framework	10,688	12,615	17.0
rustls	rustls	TLS implementation	16,866	23,658	27.3
sccache		Distributed build cache	23,202	24,673	38.3
nalgebra		Numerics library	31,951	37,143	20.8
image		Image processing library	19,718	38,165	36.4
hyper		HTTP server	15,082	39,981	50.6
rg3d		3D game engine	54,426	60,798	17.6
rav1e		Video encoder	50,294	74,667	86.4
RustPython	vm	Python interpreter	48,124	94,722	28.9
Total:			285,975	416,957	

Table 1. Dataset of crates used to evaluate slicer precision, ordered in increasing number of slices. Each project often contains many crates, so a sub-crate is specified where applicable, and the root crate is analyzed otherwise.

Appendix A.3. For each crate, we ran the slicer on every function. For a given function with n local variables, we generated n slicing criteria for the value of each variable at all return points of the function. For each criterion, we ran the modular slicer under each of the eight conditions. Finally, for each condition, we computed the size of the computed slice in the number of MIR instructions.

Given the crate-level limitation of WHOLE-PROGRAM discussed at the beginning of Section 6, we also wanted to estimate the impact of this limitation. That is, if a hypothetical tool *could* recurse into every dependency, would precision be substantively improved? We cannot answer the question directly due to the Rust compiler. However as a proxy, for each slice with the WHOLE-PROGRAM modification, we computed whether or not the slicer ever reached a crate boundary. Then we compared the slice size differences between slices that reached boundaries and those that did not.

Efficiency was not a primary concern for our evaluation because WHOLE-PROGRAM is not a truly fair comparison. WHOLE-PROGRAM is the most context-sensitive possible slicer, whereas many interprocedural slicing algorithms improve efficiency by using context-insensitive approximations that e.g. only analyze a function once by grouping together multiple call-sites to the same function. But the execution times were nonetheless small — the median execution time for the unmodified slicer was $109\mu\text{sec}$ (max 1.1sec), and for the fully modified slicer was $194\mu\text{sec}$ (max 2.9sec).

6.3 Quantitative results

We observed no meaningful patterns from the interaction of modifications, e.g. MUT-BLIND and REFERENCE-BLIND being active simultaneously versus separately. So to simplify our presentation of results, we focus only on four conditions: three for each modification individually active with the rest disabled, and one for all modifications disabled, referred to as MODULAR.

6.3.1 *WHOLE-PROGRAM*. Our first evaluation question is: how does MODULAR compare against a slicer with more information, i.e. WHOLE-PROGRAM? To answer this, Figure 5 shows the slice size differences between MODULAR and WHOLE-PROGRAM. Figure 5-left shows that the vast majority of differences are close to 0. Using a log-log scale to magnify the contours of the distribution, Figure 5-center shows that the data breaks down into two categories: $\% \text{ diff} = 0$, where the slices are exactly the same, and $\% \text{ diff} > 0$, where MODULAR is larger by some amount. A significant majority of slices are exactly the same: 95.4% of slices have $\% \text{ diff} = 0$.

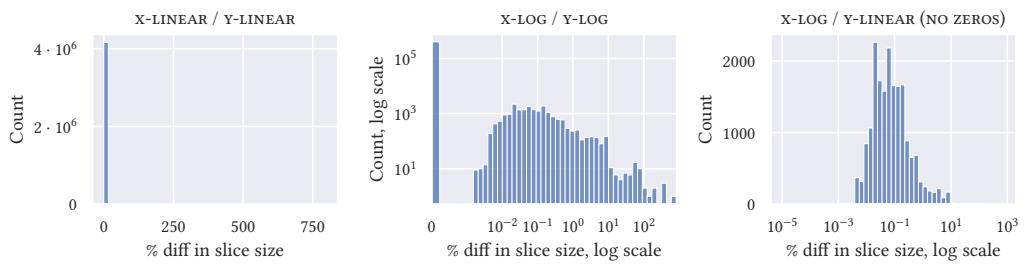


Fig. 5. Histogram of slice size differences between WHOLE-PROGRAM and MODULAR. Left: linear/linear axes, plot is mostly empty because values close to 0 dominate. Middle: log/log axes, with a bar just for zero values at the far left. Right: log/linear axes, but zero-values removed from the plotted distribution.

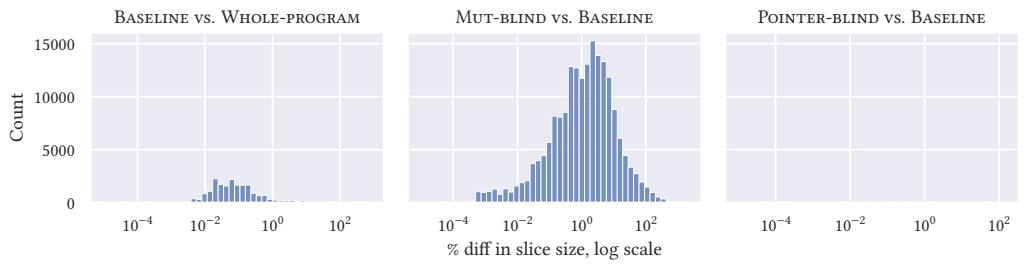


Fig. 6. Distribution of non-zero slice size differences between each modification condition and MODULAR. Plots share axes to compare relative frequency and magnitude of non-zero slices. MUT-BLIND slices very often differ from MODULAR, WHOLE-PROGRAM significantly less so, and REFERENCE-BLIND almost never (infrequently enough that they do not show up on for this shared y-axis).

Focusing on the non-zero cases, Figure 5-right shows the distribution of non-zero data on a linear y-axis. These 4.6% of slices are log-normally distributed with a geometric mean of 7.9%, lying around 10^{-1} on the log scale. These statistics indicate that the modular approximation is equally precise as inspecting definitions in nearly every slice. And when the approximation is less precise, the difference is relatively small – only 7.9% on average.

In the WHOLE-PROGRAM condition, 66.3% of slices reached a crate boundary, suggesting that the limitation does frequently affect slices. However, for boundary slices, 93.1% of them were the same size as the corresponding MODULAR slice, while for non-boundary slices, 99.5% of them were the same as MODULAR. This result suggests that exhaustive recursion didn't necessarily produce more precise slices in general compared to limited recursion.

6.3.2 MUT-BLIND and REFERENCE-BLIND. Next, we answer the question: how does MODULAR compare against a slicer with less information, i.e. MUT-BLIND and REFERENCE-BLIND? The slice size difference distributions exhibit the same zero vs. non-zero breakdown as with WHOLE-PROGRAM, so we similarly describe them by the fraction of zero cases along with the distribution of non-zero cases shown in Figure 6. MUT-BLIND increased slice sizes in 45.8% of cases for a geometric mean of 100.3% (mean around 10^0 in the Figure 6-center) in such cases. By contrast, REFERENCE-BLIND only increased slice sizes in 0.2% of cases by a geometric mean of 25.3%. REFERENCE-BLIND made so

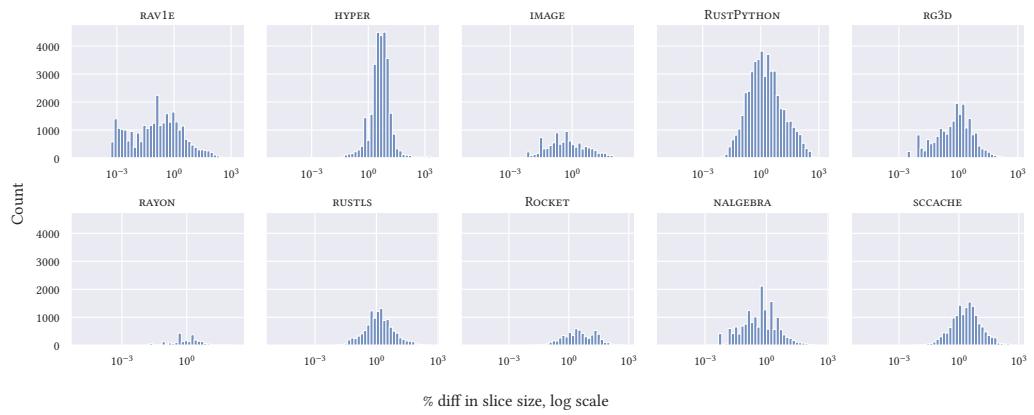


Fig. 7. Distribution of non-zero slice size differences between **MUT-BLIND** and **MODULAR** faceted by crate.

little impact that on a shared axis with the other conditions, its data is not visible in a histogram (Figure 6-right).

These results suggest that the distinction between immutable and mutable references is essential for the precision of the modular approximation. Ignoring this distinction increases slice sizes $2\times$ in nearly half of all slices! The use of lifetimes for pointer analysis was comparatively unimportant for precision, as a naive aliasing assumption only changed a fraction of a percent of slices.

6.3.3 Variation between crates. To understand the variance in slice size differences between different crates, Figure 7 shows the distribution of non-zero data for **MUT-BLIND** vs. **MODULAR** per-crane. Some crates have substantively more non-zero slices than others. To some extent, this disparity is expected because each crate has a different number of slices, e.g. **RustPython** has 94k slices while **rayon** has 10k slices.

However, **hyper** and **image** have roughly the same number of slices, but **hyper** (72.4% of slices are non-zero) is significantly more affected by **MUT-BLIND** than **image** (27.2% non-zero). Based on manual inspection of functions with non-zero slices in each crate, one explanation is simply the differing code styles between application domains. For instance, many large **image** functions are format encoder/decoders that have primarily either mutable references to buffers or owned values copied out of those buffers, but no immutable references. By contrast, many large **hyper** functions involve e.g. validating HTTP requests that primarily use shared references. Therefore **MUT-BLIND** would be more likely to affect **hyper** functions than **image** functions.

6.4 Qualitative analysis

The statistics convey a sense of how often each condition influences slice sizes. But it is equally important to understand the kind of code that leads to such differences. For each condition, we manually inspected a sample of cases with non-zero slice size differences vs. **MODULAR**.

6.4.1 Whole-program vs. modular slices. As highlighted at the beginning on Section 6, one common difference between whole-program and modular slices is functions which take a mutable pointer as input for the purposes of passing the permission off to an element of the input.

```

1030 1 fn crop<I: GenericImageView>(
1031 2     image: &mut I, x: u32, y: u32,
1032 3     width: u32, height: u32
1033 4 ) -> SubImage<&mut I> {
1034 5     let (x, y, width, height) =
1035 6         crop_dimms(image, x, y, width, height);
1036 7     SubImage::new(image, x, y, width, height)
1037 8 }

```

1038
1039 only depends on a subset of a function's inputs. A modular slice assumes all inputs are relevant to
1040 all mutations, but this is naturally not always the case.

1041 For example, this function
1042 from nalgebra returns a
1043 boolean whose value solely
1044 depends on the argument
1045 diag. However, as per the
1046 modular approximation in
1047 Principle 3, a modular slice
1048 on a call to this function
1049 would assume that all inputs
1050 are relevant to the boolean

1051 The final common case arises from unsafe code. As discussed in Section 2.4, the pattern of interior
1052 mutability uses unsafe code to hide mutations behind an immutable reference. But another kind of
1053 unsafe code came from API designers using mutable references to indicate that mutations were
1054 occurring in external C libraries.

```

1055 1 fn draw(&mut self, state: &mut State, ...) {
1056 2     let fbo: GLuint = self.id;
1057 3     state.set_framebuffer(fbo);
1058 4     unsafe { gl::DrawElements(...); }
1059 5     // many other statements elided...
1060 6 }

```

1062 draws into the previously bound framebuffer. Importantly, draw does *not* have to be **&mut self**,
1063 because the only ownership-relevant action is a read to **self.id**. However, the API designer chose
1064 to use **&mut self** because it represents the fact that the OpenGL framebuffer is being mutated.

1065 Therefore, a whole-program slice on **framebuffer.draw(..)** would consider **framebuffer** not to
1066 be mutated, and **draw** would not be part of the slice. By contrast, the modular slice would include
1067 **draw** in the slice because of the **&mut self** input type. From a strict perspective, the WHOLE-PROGRAM
1068 slice is sound and more precise, because no value *defined within Rust* is changed as a result of **draw**.
1069 However, from a system-level perspective, the MODULAR slice is correct and the WHOLE-PROGRAM
1070 slice is unsound because the value *represented* by **framebuffer**, the screen, is being mutated.

1071

1072

1073 **6.4.2 Mutability.** The reason MUT-BLIND produces many larger slices than MODULAR is quite
1074 simple — many functions take immutable references as inputs, and so many more mutations have
1075 to be assumed. Importantly, due to higher-order functions, this issue could not be solved just with
1076 a whole-program analysis.

1077

1078

For example, the function **crop** from the **image** crate returns a mutable view on a sub-image of a bigger image. No data is mutated, but rather the mutable permission is passed from whole image to sub-image. However, a modular slice on the **image** input would assume that **image** is mutated by **crop**.

Another difference arises when a value

```

1 1 fn solve_lower_triangular_with_diag_mut<R2, C2, S2>(
2     &self, b: &mut Matrix<N, R2, C2, S2>, diag: N,
3 ) -> bool {
4     if diag.is_zero() {
5         return false;
6     }
7     // function logic with no returns...
8     true
9 }

```

For example, consider this heavily simplified function from the rg3d game engine. This engine uses OpenGL to draw to the screen, and OpenGL uses a implicitly stateful API. Here, a frame buffer carries an identifier that is passed to OpenGL via **set_framebuffer**. Then **DrawElements**

via **DrawElements** draws into the previously bound framebuffer. Importantly, draw does *not* have to be **&mut self**, because the only ownership-relevant action is a read to **self.id**. However, the API designer chose to use **&mut self** because it represents the fact that the OpenGL framebuffer is being mutated.

```

1079 1 fn read_until<R, F>(io: &mut R, func: F)
1080 2   -> io::Result<Vec<u8>>
1081 3   where R: Read, F: Fn(&[u8]) -> bool
1082 4   {
1083 5     let mut buf = vec![0; 8192];
1084 6     let mut pos = 0;
1085 7     loop {
1086 8       let n = io.read(&mut buf[pos..])?;
1087 9       pos += n;
1088 10      if func(&buf[..pos]) { break; }
1089 11      // ...
1090 12    }
1091 13  }

```

1092 hypothesis is that it is simply uncommon to have multiple mutable references of the same type
 1093 within a given function, and so the lifetime-based pointer analysis is not useful most of the time.
 1094

1095 However, such cases do nonetheless arise.
 1096 For instance, in rg3d this function takes two
 1097 arguments of type `&mut` `FbxComponent`. Most of
 1098 the code modifies the parent argument. However,
 1099 a naive pointer analysis must assume that
 1100 parent and child could alias, and therefore mu-
 1101 tating parent also mutates child. Therefore a
 1102 REFERENCE-BLIND slice on child would unnec-
 1103 essarily include these mutations that would be
 1104 avoided in MODULAR.

1105 7 RELATED WORK

1106 Historically, the primary technique for backwards static slicing has been analysis of system depen-
 1107 dence graphs (SDG), introduced by Horwitz et al. [1988]. SDGs encode the call graph of a program to
 1108 enable efficient interprocedural (whole-program) static analysis. SDGs are combined with a pointer
 1109 analysis like the methods of Andersen [1994] or Steensgaard [1996] to handle indirect mutations
 1110 and function calls. For instance, the recent dg slicer for LLVM uses these techniques [Chalupa 2016].

1111 Fundamentally, prior work on slicing has been limited by the requirement to analyze existing
 1112 languages like C and Java. By focusing on newer languages with more expressive type systems,
 1113 we are able to build on three lines of research: modular static analysis, ownership types, and
 1114 information flow.

1115 7.1 Modular static analysis

1116 The key technique to modularizing static analysis is the ability to symbolically summarize a
 1117 function’s behavior. Rountev et al. [1999] and Cousot and Cousot [2002] pioneered this technique of
 1118 symbolic procedure summaries. Later work built on this framework for a variety of static analyses.
 1119 Gulwani and Tiwari [2007] create symbolic summaries for the domains of unary uninterpreted
 1120 functions and linear arithmetic. Yorsh et al. [2008] designed a framework of micro-transformers
 1121 which they applied to null-checking and typestate. Dillig et al. [2011] apply procedure summaries
 1122 to perform a modular pointer analysis on C and C++ programs.

1123 From this perspective, ownership types and their corresponding inference/checking are just
 1124 another system of symbolic procedure summaries. As discussed in Section 2.3, lifetime parameters

1125 For instance, this function from hyper repeatedly calls an input function `func` with segments of an input buffer. Without a control-flow analysis, it is impossible to know what functions `read_until` will be called with. And so MUT-BLIND must always assume that `func` could mutate `buf`. However, MODULAR can rely on the immutability of shared references and deduce that `func` could not mutate `buf`.

1126 6.4.3 *Lifetimes.* Surprisingly, information about aliases from lifetimes seemed to make almost no impact on slice sizes. One

```

1  fn link_child_with_parent_component(
2   parent: &mut FbxComponent,
3   child: &mut FbxComponent,
4   child_handle: Handle<FbxComponent>,
5 ) { match parent {
6   FbxComponent::Model(model) => {
7     model.geoms.push(child_handle),
8   },
9   // ...
10  } }

```

1128 are the key mechanism that enable symbolic expression of dataflow within a function's type
 1129 signature.

1130 7.2 Ownership types and Rust

1132 Rust and Oxide's conceptions of ownership derive from Clarke et al. [1998] and Grossman et al.
 1133 [2002]. For instance, the Cyclone language of Grossman et al. uses regions to restrict where a
 1134 pointer can point-to, and region variables to express relationships between regions in a function's
 1135 input and output. A lifetime is similar in that it annotates the types of pointers, but differs in how
 1136 it is analyzed.

1137 Matsakis [2018] first observed that Rust's lifetimes could be analyzed and interpreted as a
 1138 set of loans rather than a set of lines. Ongoing efforts in the Rust community seek to integrate
 1139 this perspective into the Rust compiler to reduce the number of valid-but-rejected programs.
 1140 Oxide [Weiss et al. 2019] recently formalized this set-of-loans model, providing the basis for our
 1141 work.

1142 Recent work have demonstrated innovative applications of Rust's type system for modular
 1143 program analysis. Astrauskas et al. [2019] embed Rust programs into a separation logic to verify
 1144 pre/post conditions about functions. Jung et al. [2020] use Rust's ownership-based guarantees to
 1145 implement provably correct program optimizations.

1146 7.3 Information flow

1147 Information flow has been historically studied in the context of security, such as ensuring low-
 1148 security users of a program cannot infer anything about its high-security internals. However,
 1149 information flow is fundamentally a generalization of dataflow to consider control-flow and effects,
 1150 both relevant to slicing. The connection between information flow and slicing was first established
 1151 in Abadi et al. [1999], where they showed that a slicing calculus could be embedded in a more
 1152 general dependency calculus (because dependency sets, like security labels, form a lattice).

1153 Hence, our work bears some resemblance to prior systems for information flow inference in
 1154 the presence of effects, such as Pottier and Simonet [2003]. One notable difference is that our
 1155 slicing algorithm has to do more work to track updates to dependency sets across effects. An
 1156 information flow system needs only to ensure e.g. for $p := e$ that p cannot be low-security when
 1157 e is high-security. By contrast, our system must propagate the dependencies of e to all possible
 1158 pointees of p .

1159 8 CONCLUSION

1160 We have demonstrated that ownership can be leveraged to build a modular slicing algorithm that
 1161 is both provably sound and reasonably precise in practice: equivalent to a whole-program slicer in
 1162 95.4% of cases, and only 7.6% larger in the remaining cases. In the short-term, these results suggest
 1163 that our modular slicer could be a practical program comprehension tool for Rust programmers. To
 1164 that end, we hope that our interactive slicer shown in Figure 4 can help users understand Rust code.

1165 More broadly, we are excited by the prospects of repurposing ownership for other forms of static
 1166 analysis. For instance, by formalizing our slicer as a general algorithm for tracking dependencies,
 1167 our work could be easily applied to other domains of information flow such as security types. Prior
 1168 work on information flow often needed to restrict the underlying language – for instance, the MLIF
 1169 language of Pottier and Simonet [2003] could not be a drop-in replacement for all existing OCaml
 1170 programs, since critical features like security-level polymorphism must be integrated manually by
 1171 the programmer.

1172 However, as discussed in Section 2.3, Rust programs already use lifetime-polymorphism to
 1173 propagate dataflow across function boundaries for modular ownership analysis. We were able

1174

1177 to implement our modular slicer without any changes to the Rust language or to existing Rust
 1178 programs. And unlike analyzing C or C++, there are no unpredictable expressiveness cliffs (as the
 1179 cliffs are explicitly demarcated by `unsafe`). By restricting the space of permissible programs with
 1180 ownership-safety, we simultaneously expand the space of program analyses that were previously
 1181 intractable in more expressive languages.

1182 REFERENCES

- 1184 Martín Abadi, Anindya Banerjee, Nevin Heintze, and Jon G Riecke. 1999. A core calculus of dependency. In *Proceedings of
 1185 the 26th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. 147–160.
- 1186 Lars Ole Andersen. 1994. *Program analysis and specialization for the C programming language*. Ph.D. Dissertation. Citeseer.
- 1187 Vytautas Astrauskas, Christoph Matheja, Federico Poli, Peter Müller, and Alexander J Summers. 2020. How do programmers
 1188 use unsafe rust? *Proceedings of the ACM on Programming Languages* 4, OOPSLA (2020), 1–27.
- 1189 Vytautas Astrauskas, Peter Müller, Federico Poli, and Alexander J Summers. 2019. Leveraging Rust types for modular
 1190 specification and verification. *Proceedings of the ACM on Programming Languages* 3, OOPSLA (2019), 1–30.
- 1191 Gogul Balakrishnan, Radu Gruian, Thomas Reps, and Tim Teitelbaum. 2005. Codesurfer/x86—a platform for analyzing x86
 1192 executables. In *International Conference on Compiler Construction*. Springer, 250–254.
- 1193 Marek Chalupa. 2016. *Slicing of LLVM bitcode*. Master's thesis. Masaryk University.
- 1194 David G Clarke, John M Potter, and James Noble. 1998. Ownership types for flexible alias protection. In *Proceedings of the
 1195 13th ACM SIGPLAN conference on Object-oriented programming, systems, languages, and applications*. 48–64.
- 1196 Keith D Cooper, Timothy J Harvey, and Ken Kennedy. 2001. A simple, fast dominance algorithm. *Software Practice &
 1197 Experience* 4, 1-10 (2001), 1–8.
- 1198 Patrick Cousot and Radhia Cousot. 2002. Modular static program analysis. In *International Conference on Compiler
 1199 Construction*. Springer, 159–179.
- 1200 Pascal Cuq, Florent Kirchner, Nikolai Kosmatov, Virgile Prevosto, Julien Signoles, and Boris Yakobowski. 2012. Frama-c. In
 1201 *International conference on software engineering and formal methods*. Springer, 233–247.
- 1202 Ron Cytron, Jeanne Ferrante, Barry K Rosen, Mark N Wegman, and F Kenneth Zadeck. 1989. An efficient method of
 1203 computing static single assignment form. In *Proceedings of the 16th ACM SIGPLAN-SIGACT symposium on Principles of
 1204 programming languages*. 25–35.
- 1205 Al Danial. 2021. cloc: Count Lines of Code. <https://github.com/AlDanial/cloc>
- 1206 Isil Dillig, Thomas Dillig, Alex Aiken, and Mooly Sagiv. 2011. Precise and compact modular procedure summaries for heap
 1207 manipulating programs. *ACM SIGPLAN Notices* 46, 6 (2011), 567–577.
- 1208 Jeanne Ferrante, Karl J Ottenstein, and Joe D Warren. 1987. The program dependence graph and its use in optimization.
 1209 *ACM Transactions on Programming Languages and Systems (TOPLAS)* 9, 3 (1987), 319–349.
- 1210 Jean-Yves Girard. 1987. Linear logic. *Theoretical computer science* 50, 1 (1987), 1–101.
- 1211 Dan Grossman, Greg Morrisett, Trevor Jim, Michael Hicks, Yanling Wang, and James Cheney. 2002. Region-based memory
 1212 management in Cyclone. In *Proceedings of the ACM SIGPLAN 2002 Conference on Programming language design and
 1213 implementation*. 282–293.
- 1214 Sumit Gulwani and Ashish Tiwari. 2007. Computing procedure summaries for interprocedural analysis. In *European
 1215 Symposium on Programming*. Springer, 253–267.
- 1216 Susan Horwitz, Thomas Reps, and David Binkley. 1988. Interprocedural slicing using dependence graphs. In *Proceedings of
 1217 the ACM SIGPLAN 1988 conference on Programming Language design and Implementation*. 35–46.
- 1218 Ganeshan Jayaraman, Venkatesh Prasad Ranganath, and John Hatcliff. 2005. Kaveri: Delivering the indus java program
 1219 slicer to eclipse. In *International Conference on Fundamental Approaches to Software Engineering*. Springer, 269–272.
- 1220 Ralf Jung, Hoang-Hai Dang, Jeehoon Kang, and Derek Dreyer. 2020. Stacked borrows: an aliasing model for Rust. *Proceedings
 1221 of the ACM on Programming Languages* 4, POPL (2020), 1–32.
- 1222 Niko Matsakis. 2018. An alias-based formulation of the borrow checker. <http://smallcultfollowing.com/babysteps/blog/2018/04/27/an-alias-based-formulation-of-the-borrow-checker>
- 1223 Matthew Might, Yannis Smaragdakis, and David Van Horn. 2010. Resolving and exploiting the k-CFA paradox: illuminating
 1224 functional vs. object-oriented program analysis. In *Proceedings of the 31st ACM SIGPLAN Conference on Programming
 1225 Language Design and Implementation*. 305–315.
- 1226 Chris Parnin and Alessandro Orso. 2011. Are automated debugging techniques actually helping programmers?. In *Proceedings
 1227 of the 2011 international symposium on software testing and analysis*. 199–209.
- 1228 François Pottier and Vincent Simonet. 2003. Information flow inference for ML. *ACM Transactions on Programming
 1229 Languages and Systems (TOPLAS)* 25, 1 (2003), 117–158.
- 1230 Atanas Rountev, Barbara G Ryder, and William Landi. 1999. Data-flow analysis of program fragments. In *Software
 1231 Engineering—ESEC/FSE’99*. Springer, 235–252.

- 1226 Josep Silva. 2012. A vocabulary of program slicing-based techniques. *ACM computing surveys (CSUR)* 44, 3 (2012), 1–41.
1227 Yannis Smaragdakis and George Balatsouras. 2015. Pointer analysis. *Foundations and Trends in Programming Languages* 2, 1
1228 (2015), 1–69.
1229 Bjarne Steensgaard. 1996. Points-to analysis in almost linear time. In *Proceedings of the 23rd ACM SIGPLAN-SIGACT
symposium on Principles of programming languages*. 32–41.
1230 Mads Tofte and Jean-Pierre Talpin. 1997. Region-based memory management. *Information and computation* 132, 2 (1997),
1231 109–176.
1232 Mark Weiser. 1982. Programmers use slices when debugging. *Commun. ACM* 25, 7 (1982), 446–452.
1233 Mark Weiser. 1984. Program slicing. *IEEE Transactions on software engineering* 4 (1984), 352–357.
1234 Aaron Weiss, Olek Gierczak, Daniel Patterson, Nicholas D Matsakis, and Amal Ahmed. 2019. Oxide: The essence of rust.
arXiv preprint arXiv:1903.00982 (2019).
1235 Baowen Xu, Ju Qian, Xiaofang Zhang, Zhongqiang Wu, and Lin Chen. 2005. A brief survey of program slicing. *ACM
SIGSOFT Software Engineering Notes* 30, 2 (2005), 1–36.
1237 Greta Yorsh, Eran Yahav, and Satish Chandra. 2008. Generating precise and concise procedure summaries. In *Proceedings of
the 35th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. 221–234.
1238 Jisheng Zhao, Michael G Burke, and Vivek Sarkar. 2018. Parallel sparse flow-sensitive points-to analysis. In *Proceedings of
the 27th International Conference on Compiler Construction*. 59–70.
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274